

Parallel and Distributed Systems - Exercise 3

CUDA - ISING MODEL

Authors:

Christoforidis Savvas, AEM: 9147, email: schristofo@ece.auth.gr

Portokalidis Stavros, AEM: 9334, email: stavport@ece.auth.gr

Code: <https://github.com/schristofo/CUDA-ising-model>

v0. Sequential: The first version of our program, basically, implements the *ising()* function given to us by the online tester. Analytically, using two arrays as requested we copy the data from one to the other, so as to use the one of them as an input and the other one as the output. This is achieved by swapping their pointers after each calculation of the new spin values. There is also a check in case of having the spin lattice unchanged using a boolean variable. More deeply, the calculation of new spin values consists of the calculation of the influence and the sign function of each influence. The influence calculation uses mod arithmetic to avoid if statements for the boundaries. Finally, in order to take into account possible floating point errors and check if the influence value is zero, we use a threshold value(10^{-7}).

v1. GPU with one thread per moment

In this part of the exercise, we were asked to modify the sequential code (v0), make a kernel function and call the kernel with a grid that matches the Ising model and one thread per moment. First, we transfer the data to GPU memory. After that, we define `<< c = blockIdx.x * blockDim.x + threadIdx.x, r = blockIdx.y * blockDim.y + threadIdx.y; >>` to calculate the weighted influence of each point's neighbors. Finally, we compute the spin of each point, according to Ising model method. Every GPU's thread will compute the spin of one point.

v2. GPU with one thread computing a block of moments

This version of the method is slightly different from the previous one, due to the fact that each thread is computing a block of moments, increasing by grid value each time. Hence, we use a grid-stride loop by moving the block $\text{gridDim.x} * \text{blockDim.x}$ in terms of X and $\text{gridDim.y} * \text{blockDim.y}$ in terms of Y. In this way, we are able to make good use of every utilizable thread.

v3. GPU with multiple thread sharing common input moments

In the last version of the function *ising()* we modified version v2 by adding two attached-to-device memory arrays for faster access. In other words we used two shared memory blocks, one for storing the *w* matrix and one for storing a block of moments and its radius, accessible from every thread in each block. The result was amazing, since time performance improved drastically.

Time Performance and Visualization:

System Characteristics / GPU model:

Intel® Core™ i7-7500U CPU @ 2.70GHz × 4 / NVIDIA GK208 (GeForce 920M)

CUDA Version 10.2.89

Optimal block size used for testing: **16 x 16**

Results:

A) Results with constant n


| n | k | V0 | V1 | V2 | V3 |
|------|-----|----------|-------|-------|-------|
| 1000 | 1 | 0.070918 | 0.062 | 0.059 | 0.053 |
| 1000 | 50 | 1.448755 | 0.751 | 0.728 | 0.550 |
| 1000 | 100 | 2.896730 | 1.445 | 1.395 | 1.063 |

| n | k | V0 | V1 | V2 | V3 |
|------|-----|-----------|--------|--------|--------|
| 5000 | 1 | 1.02888 | 0.503 | 0.475 | 0.394 |
| 5000 | 50 | 35.714700 | 11.264 | 10.797 | 6.753 |
| 5000 | 100 | 79.232514 | 22.883 | 20.225 | 14.787 |


| n | k | V0 | V1 | V2 | V3 |
|-------|-----|------------|---------|---------|--------|
| 10000 | 1 | 4.457116 | 1.812 | 1.938 | 1.459 |
| 10000 | 50 | 155.13 | 60.418 | 58.262 | 39.293 |
| 10000 | 100 | 388.272135 | 132.711 | 122.130 | 90.952 |

Graphs:

Execution time [N = 5000]



Execution time [N = 10000]





B) Results with constant k

| n | k | V0 | V1 | V2 | V3 |
|-------|----|-----------|--------|--------|--------|
| 1000 | 25 | 0.763684 | 0.389 | 0.3844 | 0.310 |
| 2000 | 25 | 6.044902 | 1.476 | 1.426 | 1.049 |
| 5000 | 25 | 23.2494 | 9.193 | 8.451 | 6.477 |
| 7000 | 25 | 45.163928 | 16.805 | 14.361 | 10.415 |
| 10000 | 25 | 85.8424 | 27.165 | 25.561 | 18.182 |

| n | k | V0 | V1 | V2 | V3 |
|-------|----|-----------|--------|--------|--------|
| 1000 | 50 | 1.422374 | 0.755 | 0.745 | 0.572 |
| 2000 | 50 | 9.992142 | 2.787 | 2.747 | 2.055 |
| 5000 | 50 | 35.9537 | 11.996 | 10.713 | 6.485 |
| 7000 | 50 | 70.163928 | 28.865 | 26.319 | 20.033 |
| 10000 | 50 | 155.12897 | 60.460 | 58.518 | 39.469 |

Graphs:



Easily, we can observe that execution time grows linearly for constant n and growing k. On the other hand, when we keep k constant and we increase n there is an exponential growth in execution time. Finally, as far as the version execution times are concerned, we see that $tv_0 > tv_1 > tv_2 > tv_3$.